

Comment on “Tainted evidence: cosmological model selection versus fitting”, by Eric V. Linder and Ramon Miquel (astro-ph/0702542v2)

Andrew R. Liddle,^{1,2} Pier Stefano Corasaniti,³ Martin Kunz,⁴
Pia Mukherjee,¹ David Parkinson,¹ and Roberto Trotta⁵

¹*Astronomy Centre, University of Sussex, Brighton BN1 9QH, United Kingdom*

²*Institute for Astronomy, University of Hawai‘i, 2680 Woodlawn Drive, Honolulu, HI 96822, USA*

³*LUTH, CNRS UMR 8102, Observatoire de Paris-Meudon,*

5 Place Jules Janssen, 92195 Meudon Cedex, France

⁴*Département de Physique Théorique, Université de Genève,*

24 quai Ernest Ansermet, CH-1211 Genève 4, Switzerland

⁵*Oxford University, Astrophysics, Denys Wilkinson Building, Keble Road, OX1 3RH, United Kingdom*

(Dated: February 4, 2008)

In astro-ph/0702542, Linder and Miquel seek to criticize the use of Bayesian model selection for data analysis and for survey forecasting and design. Their discussion is based on three serious misunderstandings of the conceptual underpinnings and application of model-level Bayesian inference, which invalidate all their main conclusions. Their paper includes numerous further inaccuracies, including an erroneous calculation of the Bayesian Information Criterion. Here we seek to set the record straight.

PACS numbers: 98.80.-k

astro-ph/0703285

I. INTRODUCTION

In a recent paper, Linder and Miquel [1] have mounted a vigorous attack on the use of model selection techniques in cosmology, particularly with regard to interpreting (forecasting) the outcome of (upcoming) surveys and in survey design applications. They instead advocate a frequentist parameter-fitting technique.

Our aim in this short note is to highlight important misunderstandings that invalidate all the main conclusions of their paper. In the process, we give a brief self-contained discussion of the model selection framework; for more details see e.g. Refs. [2, 3, 4]. In the Appendix we highlight some specific inaccuracies in Ref. [1], many of which are consequences of the general misunderstandings outlined in the main body of this Comment.

II. WHAT IS BAYESIAN MODEL SELECTION?

In Bayesian inference, model parameters are taken as random variables, because this allows propagation of the experimental measurement errors into self-consistent probabilistic statements about parameter uncertainties.

The first step of Bayesian parameter estimation is the choice of a model (M_i), which specifies a set of parameters ($\vec{\theta}_i$) to be varied in fitting to the data, along with a set of prior probability ranges $P(\vec{\theta}_i|M_i)$ for those parameters. Given a particular set of data D , the likelihood $P(D|\vec{\theta}_i, M_i)$ is used to update the prior probabilities to the posterior

$$P(\vec{\theta}_i|D, M_i) = \frac{P(D|\vec{\theta}_i, M_i) P(\vec{\theta}_i|M_i)}{P(D|M_i)}. \quad (1)$$

The posterior $P(\vec{\theta}_i|D, M_i)$ contains all the information

about the state of knowledge on the parameters $\vec{\theta}_i$ after the arrival of the data. From this one can construct ‘credible intervals’, i.e. ranges encompassing say 68% or 95% of posterior probability for the parameters.

We remark that Bayesian credible intervals have a profoundly different meaning from frequentist confidence regions, where model parameters are not random variables but fixed unknown quantities. The fact that the two intervals are formally equal in the case of a Gaussian likelihood (and flat priors, in the Bayesian scheme) is traceable to the symmetry between the measured mean and the ‘true’ mean entering the Gaussian distribution. This formal equivalence can engender considerable confusion as to the different interpretations of the final result (for a detailed discussion see Ref. [5]).

Bayesian model selection (or comparison) is the extension of the parameter estimation framework to include multiple models, with different parameter vectors and priors. Bayes theorem can be applied again to update a prior model probability by the *evidence*, also known as the marginal likelihood of the model, which is the normalization constant in Eq. (1)

$$P(D|M_i) = \int P(D|\vec{\theta}_i, M_i) P(\vec{\theta}_i|M_i) d\vec{\theta}_i. \quad (2)$$

The evidence is the probability of the data given the model. Bayes theorem is then used to obtain the probability of the model given the data,

$$P(M_i|D) \propto P(D|M_i) P(M_i), \quad (3)$$

where $P(M_i)$ is the prior model probability. It is clear from the above equations that the evidence is the basis of the model comparison and is built upon the parameter estimation step.

All of the above is uncontroversial mathematics, providing a consistent and systematic inference system for

evolving probabilities in light of experimental data. Any controversy about Bayesian methods centres around the explicit need to state the full set of prior information in order to do any calculation. The framework provides no guidance as to how to do this; instead physical insight is needed to select suitable models for comparison with data, and to assess their initial probability and the priors on the model parameters in advance of that comparison. Within a model, prior parameter ranges can be thought of as plausible regions of parameter space that are accessible to the model.

Bayes theorem can be thought of as a decomposition of the final result into prior knowledge and the likelihoods measuring the information coming from the data. Indeed, the Bayesian formalism forces us to state explicitly which part of the result is due to our assumptions, and which part is driven by the data. The hope or expectation, both at the parameter level and model level, is that data will be obtained of sufficient quality to overturn incorrect prior hypotheses. If there is a broad range of possible models, corresponding to different prior choices, it will require more data to converge to a robust conclusion. But in the Bayesian approach data will eventually overcome prior choices; the wider the range of plausible priors, the more data we can expect to need before a firm conclusion can be drawn.

The Bayesian evidence sets up a tension between the ability of a model to fit the data and the prior predictiveness of the model, in a quantitative implementation of Occam’s razor. Note that that we prefer to use the term ‘predictiveness’ rather than ‘simplicity/complexity’; the former is what is actually rewarded by the evidence, and is not necessarily directly related to, for instance, the number of parameters. The models that do best are the ones that make specific predictions that later turn out to fit the data well. Less predictive models, even if they can fit the data as well, score more poorly. The Bayesian evidence has been widely applied to cosmological problems in recent years [3, 4, 6, 7, 8, 9, 10, 11, 12].

Some statistics have been extensively used as proxies to the actual evidence, such as the Bayesian Information Criterion (BIC) [13, 14]. But unlike the evidence these approximations are often biased, and by construction disfavour models with more parameters, even when those parameters are not constrained by the data (see IIC for more on this in relation to the evidence, and Ref. [14] for a discussion of the limitations of some information criterion based approaches). Wherever possible the full evidence should be used.

III. MISCONCEPTIONS ABOUT MODEL SELECTION

The paper of Linder and Miquel [1] launches a primarily rhetorical attack on the model selection framework. We will argue here that the paper contains numerous factually-incorrect statements. These appear largely

to be traceable to three fundamental misunderstandings concerning the Bayesian framework and its applications, which we now describe.

A. Model selection does not replace parameter estimation. It extends it.

Linder and Miquel appear to believe that model selection and parameter estimation are competing techniques. This is incorrect. As described above, model selection extends the Bayesian framework to the model level. Within each model, parameter estimation is carried out in the usual manner. This would include, as usual, goodness-of-fit and data subset consistency checks.

Specifically, we see that parameter estimation corresponds to model selection where the prior model probabilities of all but one model have been set to zero. This seems a regressive step; one can hardly claim that our understanding of, for instance, dark energy is so good that we should focus on only one possible description.

From this perspective, the need to choose model priors is clearly an advantage, not a drawback. Parameter estimation corresponds to one particular choice of those priors. By acknowledging that other choices are possible, a much more wide-ranging and robust investigation of the possible outcomes of future experiments can be made, as was done in Ref. [11].

A further advantage of model selection is that it allows one to ask new types of question. As it subsumes parameter estimation, one can obviously still ask about parameter confidence ranges, for instance, either model-by-model or via Bayesian model averaging as in Ref. [11]. But one can also ask whether entire models are excluded by data at a given strength of evidence, based on their posterior model probability, or whether data provide support for additional model parameters. Indeed, the current leading questions in dark energy studies are of model selection type, viz. is the equation of state w equal to -1 or is it variable? In the latter case, is w constant or time-varying? One can also compare models that are not nested, for instance is quintessence a better description of the data than a modified gravity model? Such questions are not accessible to parameter fitting analyses and often cannot even be phrased in frequentist terms.

An important application of model selection is survey forecasting and design, where one assesses or constructs a survey in order to optimize the ability to answer a particular question or questions [4, 8, 12]. The details will inevitably depend to some extent on prior assignments, and it is of course important to vary these within reasonable ranges. Model selection forecasting allows optimization for a broader range of possible questions.

Linder and Miquel also claim that survey design based on model selection is betting on the absence of structure in the possible parameter space, apparently confusing the space of possible ‘true’ models with the likelihood in that space given a particular ‘true’ model. The opposite is

true. By including several models, one can focus attention on particular regions of parameter space that are especially well motivated, for instance Λ CDM, or the locations predicted by one-parameter quintessence models. In fact it is parameter estimation that assumes that the parameter space is a blank canvas in which each point is of equal value.

B. Physical intuition and priors are the same thing!

Linder and Miquel criticize the Bayesian methodology for giving results that are often dependent on prior assumptions, and simultaneously claim that it seeks to avoid, or even prevent, use of physical intuition. Apparently they have not realised that physical intuition and priors are the same thing! After all, where do the models come from that we decide to compare to the data? What decides their prior model probabilities, and the reasonable ranges for their parameters? This is where the physics comes in. The mere fact that it can be difficult to put our physical intuition in quantitative terms by selecting prior ranges and prior probabilities is no good reason to give up the exercise.

The Bayesian model selection framework, by allowing us to specify multiple models with both model and parameter priors, *maximizes* our chance to incorporate physical intuition into data analysis. Linder and Miquel’s claims to the contrary hold no substance at all. From this perspective, the prior dependence in Bayesian analysis should be viewed in a positive light, not a negative one, as it allows different intuitions to be tested. Bayes’ theorem provides a convenient decomposition into the parts of the conclusions that are data-driven (the parameter and model likelihoods) and those that are prior-driven (the physical intuition), and so one can always keep track of the balance between those two. As long as the data cannot decide the issue, our physical intuition influences the outcome of our conclusions, but in the Bayesian framework we are made explicitly aware of this situation through the need to specify an explicit prior. There is no inference without assumptions. As the amount and quality of data increases, the priors become less important and the conclusions based on our expectations are replaced by conclusions based on actual data. This is how physics should work.

C. Model selection does not act against models whose parameters cannot yet be measured.

Linder and Miquel give a historical overview, titled ‘reality check’, which seeks to show by example that model selection techniques, if applied in the past, would have led researchers astray. In our view all of this section is incorrect, as we explain in detail in the Appendix. Here we will address the reason why Linder and Miquel have gone astray.

Their main mistake is a failure to recognize the difference between two distinct circumstances. The first is a situation where a phenomenon could have been discovered, but wasn’t; this corresponds to a likelihood function well localized within the prior of the relevant parameter, but consistent with a zero value. Model selection statistics act against models with the extra parameter in that case (an example being spatial curvature). The second is the situation where observations were of insufficient power to constrain the parameter, corresponding to a flat or nearly flat likelihood across the prior. In this case, the contribution of the parameter factorizes out of the evidence integral, leaving it unchanged. Therefore *Bayesian model selection does not act against parameters that are unconstrained by existing data* (see Ref. [9] for a detailed discussion). Comparisons of such models are inconclusive, awaiting new data. All the examples they give purporting to show model selection going astray are actually in the second category, and not in the first as they say.

Their misunderstanding can be partially traced to their use of the BIC [13]. This model selection criterion assumes that all parameters are well measured. If this is not the case, then the BIC will exclude models that are perfectly acceptable when using full Bayesian model comparison, as e.g. demonstrated in Ref. [7]. There it was shown that the BIC rules out the “kink” parametrisation of the dark energy equation of state, in disagreement with the full Bayesian evidence. Indeed, in the derivation of the BIC only the scaling with the number of data points N was kept, while even in the idealized case of linear models with Gaussian errors the overall scaling of the log-evidence for a model with k parameters is rather $k \ln(N/k)$ [9]. Additionally there is a term that depends on the size of the error bars relative to the size of the prior, which often dominates. For these reasons, the BIC tends to give an unrealistically high penalty to extra parameters, compared to the full Bayesian evidence, if its underlying assumptions are not met. Only the evidence is a full implementation of Bayesian model comparison.

IV. CONCLUSIONS

We regard model selection techniques as a powerful tool for cosmologists, both for data analysis and for survey forecasting and design. They broaden the range of questions one can ask of present and future observations, and can be applied in a consistent and rigorous framework. While there remains room for debate about the relative merits of frequentist and Bayesian approaches in cosmology, we believe that the many demonstrable flaws of the Linder–Miquel paper prevent it from contributing constructively to that debate.

Acknowledgments

P.S.C. was supported by CNRS with contract No. 06/311/SG, M.K. by the Swiss NSF, P.M. and D.P. by PPARC (UK), and R.T. by the Royal Astronomical Society’s Sir Norman Lockyer Fellowship and by St. Anne’s College, Oxford. We thank Eric Linder and Ramon Miquel for discussing the content of their paper with us. A.R.L. additionally thanks Gang Chen, Mark Neyrinck, Adrian Pope, and Istvan Szapudi for discussions.

APPENDIX A: DETAILED CRITIQUE

In this Appendix we provided a detailed critique of some of points made by Linder and Miquel.

1. Linder and Miquel: Section III/V

In our view, the discussion in their section III, seeking instances from history where model selection would have misinformed, is entirely wrong or irrelevant. Since this is the sole motivation for their Section V, it too has no validity. We take their account paragraph by paragraph.

a) This paragraph claims that, pre-1998, model selection would have dismissed the now-favoured Λ CDM model. Absolutely not! Data before that epoch were unable to meaningfully constrain Λ . As discussed above, the comparison would have been inconclusive. This is in accord with the fact that in the 1990s papers typically considered several cosmological models, including Λ CDM, on a roughly equal footing. In 1998 better data came along able to rule out the critical-density and open models, at which point model selection would correctly pick out the dark energy model.

b) This paragraph mentions Feynman’s repackaging of all equations of nature into $\bar{U} = 0$. We can see no relevance in this point. Repackaging equations does not change the number of model fit parameters, and hence affects neither parameter estimation nor model selection.

c) This paragraph claims that before 1992 model selection would have argued against structure in the cosmic microwave background (CMB), on the grounds that $C_\ell = 0$ is simpler than independently specifying each C_ℓ . Absolutely not! This point confuses the data and the models. No one has ever thought that a separate specification of each C_ℓ was a model, and certainly not in 1992. Indeed, clearly the acceptable models of the time, the CDM family with or without Λ , all predicted CMB structures that were indeed subsequently seen. Not that this has anything much to do with model selection; models cannot be rejected before obtaining data that actually constrains them.

d) This paragraph notes that the galaxy two-point correlation function was for many years thought to be a power-law, without an underpinning physical model.

However, that the power-law model is no longer considered is irrelevant. Would any more physical model have been wrongly ruled out by model selection, had they existed? No is the answer. When data improved, would model selection support them over the power-law model? Yes. As it should.

e) This paragraph makes a point about there being the deeper physics of the halo model behind the matter power spectrum, but this has nothing to do with cosmological model selection.

f) This paragraph claims that the modern electro-weak theory would have been rejected by model selection had it existed contemporaneously with the Fermi theory during its heyday. Absolutely not! Had the Glashow–Weinberg–Salam model been around in say 1950, it would *not* have been ruled out by model selection because all of its parameters were poorly constrained. Model selection would have been unable to distinguish it from the simpler Fermi model. Later on better data came along and ruled out the simpler model. Just as it should.

g) This paragraph makes a point that seemingly complex phenomenology may have a simple underlying structure, e.g. atomic spectra. There is some relevance to this point, though a ‘complicated’ say two-parameter equation of state model for dark energy is unlikely to have a substantially simpler ‘fundamental’ description. Nevertheless, if improved physical understanding comes along and creates a compelling model of that type, then that is the time to try out model selection statistics on it. Such a model can hardly be ruled out before it even exists, nor tested until its predictions are defined.

2. Linder and Miquel: Section IV

In Section IV the authors advocate a frequentist ‘rejection of null hypothesis’ test where Λ CDM is the null hypothesis. This is done by simulating data only for Λ CDM and drawing likelihood contours, with the viability of Λ CDM then interpreted according to the position of actual measurements with respect to those contours. Note that this approach seeks to rule out Λ CDM in favour of a more general dark energy model without ever computing the probability of the data under the latter model.

We first note that the quantity they compute and call BIC is *not* the BIC. A giveaway is that they are claiming that the lower likelihood models are preferred. The correct computation of the BIC requires simulation of data at each point in the parameter space, and then a model comparison test of Λ CDM versus the two-parameter dark energy model at each point. We carried out exactly such an analysis, computing the full evidence rather than the BIC, in Ref. [8]. Linder and Miquel only simulate Λ CDM, and then simply flip the sign of the relative log-likelihood, which is not equivalent.

On top of the above flaw, Linder and Miquel’s argument then goes on to describe a situation in which a frequentist analysis delivers a 90% confidence contour

around the Λ CDM model (based on synthetic data) in the (w_0, w_a) plane, and claims that a measurement lying outside that region would exclude the Λ CDM at the 90% confidence level. This is an incorrect statement in frequentist statistics, as it slips in the wrong assumption that the probability of the data given the hypothesis (i.e., the frequentist confidence region) is the same as the probability of the hypothesis given the data.¹ The latter quantity is undefined for a frequentist, for whom a hypothesis is either true or false (although we might not know which one is true) and a probabilistic statement about it would be meaningless. For a Bayesian of course the two are related through Bayes theorem.

Linder and Miquel also voice their discontent about the BIC condition being stronger, i.e. making it harder to rule out Λ CDM. But this is hardly surprising, being just a manifestation of Lindley’s well-known ‘paradox’ [15]; as summarized in Appendix A of Ref. [4] (see Figure A1), in general frequentist significance tests do not agree with Bayesian model selection, since the former ignore the information gained through the data. This is evident in Figures 1 and 2 of Ref. [8], which is exactly the comparison Linder and Miquel are trying to make. There is no basis to claim, using a frequentist significance test, that the BIC ‘spuriously rules out’ a particular set of models, because there is no basis to take the frequentist result as the ‘truth’. We could equally well say that model-level Bayesian inference demonstrates that parameter estimation ‘spuriously rules out Λ CDM’ in those circumstances.

It is true that model selection gives a larger parameter area in which Λ CDM would not be ruled out even if it is wrong, though usually returning an inconclusive verdict in that case, to be deferred to future data. The trade-off is that parameter estimation techniques applied for model comparison are much more likely to rule out Λ CDM even if it is right (by not recognizing that the data has lower probability under a less predictive model). One cannot win on both sides of that coin.

Independently of the above misconceptions, Linder and Miquel further claim that if Λ CDM results in a poor likelihood in light of new data, then it should be rejected in favour of a more general (less predictive) model, i.e. one in which w_0 and w_a vary freely over any range. It is not clear for example why *this* model was chosen instead of, for example, one where $w(z)$ varies in 1000 redshift bins, which would probably achieve an even better fit. In the Bayesian framework we can admit all those models, assigning a prior probability to each that reflects our relative degree of belief based on our understanding of the

physical processes at work. One then goes on to compute the posterior probabilities for each of the models.

3. Linder and Miquel: Section II

We disagree with all statements in Section II of their paper implying that parameter estimation and model selection are distinct endeavours. In addition we note

1) The statement that we wouldn’t want to throw away a tree containing one fit fruit is misleading. If the fit fruit is a better fit than those on other trees, then the goodness-of-fit will be rewarded by model selection. If it is no fitter than those on smaller trees (and is everywhere constrained meaningfully by data) then of course we *do* want to throw away that tree: this is what Occam’s Razor is all about and without it we have no control over arbitrarily complex models.

2) It is implied that model selection might disadvantage fundamental models that might have apparently complicated phenomenological manifestations. Specific examples mentioned are braneworld models of modified gravity and inverse power-law potentials. This criticism is not true at all; people are welcome, indeed encouraged, to deploy fundamental parameters in model selection rather than phenomenological ones where possible.

4. Rhetoric

We end by pointing out that this paper uses the rhetorical trick of attributing, without citation, and then rebutting, some vaguely ridiculous assertions supposedly held by model selection advocates. For instance, no-one has suggested that model selection techniques should be ‘blindly applied’ without regard to physical insight, and if they had it would have been a pretty ludicrous suggestion. No one has claimed that parameter fitting is ‘misguided’, it being a key part of the inference procedure, though we have indeed argued that it is inadequate if one wishes to answer questions phrased at the model level (e.g. is quintessence a better description of data than a particular modified gravity model). We are also unaware of any cases where ‘overenthusiastic application of model selection led to some claims about the probability of future experiments failing to see characteristics such as dynamics that current data cannot access’, though we may have been enthusiastic about being able to make probabilistic forecasts under carefully-defined prior assumptions [4, 10, 11, 12].

¹ To convince oneself of the difference between the two quantities, imagine selecting a person at random — the person can either be male or female (our hypothesis). If the person is female, her probability of being pregnant (our data) is about 3%, i.e. $P(\text{pregnant}|\text{female}) = 0.03$. However, if the person is pregnant, her probability of being female is much larger than that, i.e. $P(\text{female}|\text{pregnant}) \gg 0.03$. For further details, see Ref. [16].

-
- [1] E. V. Linder and R. Miquel, astro-ph/0702542v2.
 - [2] H. Jeffreys, *Theory of Probability*, 3rd ed, Oxford University Press (1961); D. J. C. MacKay, *Information theory, inference and learning algorithms*, Cambridge University Press (2003); P. Gregory, *Bayesian Logical Data Analysis for the Physical Sciences*, Cambridge University Press (2005).
 - [3] A. R. Liddle, P. Mukherjee, and D. Parkinson, *A&G* **47**, 4.30 (2006), astro-ph/0608184; A. Lasenby and M. P. Hobson, *Proc. Sci. (CMB2006)* 014.
 - [4] R. Trotta, astro-ph/0504022.
 - [5] R. D. Cousins, *Am. J. Phys.* **63**, 5 (1995).
 - [6] A. Jaffe, *Astrophys. J.* **471**, 24 (1996), astro-ph/9501070; P. S. Drell, T. J. Loredo, and I. Wasserman, *Astrophys. J.* **530**, 593 (2000), astro-ph/9905027; M. V. John and J. V. Narlikar, *Phys. Rev. D* **65**, 043506 (2002), astro-ph/0111122; A. Slosar et al., *Mon. Not. Roy. Astron. Soc.* **341**, L29 (2003), astro-ph/0212497; T. D. Saini, J. Weller, and S. L. Bridle, *Mon. Not. Roy. Astron. Soc.* **348**, 603 (2004), astro-ph/0305526; P. J. Marshall, M. P. Hobson, and A. Slosar, *Mon. Not. Roy. Astron. Soc.* **346**, 489 (2003), astro-ph/0307098; A. Niarchou, A. Jaffe, and L. Pogosian, *Phys. Rev. D* **69**, 063515 (2004), astro-ph/0308461; G. Lazarides, R. Ruiz de Austri and R. Trotta *Phys. Rev. D* **70** 123527 (2004), hep-ph/0409335; M. Beltrán, J. Garcia-Bellido, J. Lesgourgues, A. R. Liddle, and A. Slosar, *Phys. Rev. D* **71**, 063532 (2005), astro-ph/0501477; P. Mukherjee, D. Parkinson, and A. R. Liddle, *Astrophys. J. Lett.* **638**, L51 (2006), astro-ph/0508461; M. Kunz, N. Aghanim, L. Cayon, O. Forni, A. Riazuelo, and J. P. Uzan, *Phys. Rev. D* **73** (2006) 023511, astro-ph/0510164; M. Bridges, A. N. Lasenby, and M. P. Hobson, *Mon. Not. Roy. Astron. Soc.* **369**, 1123 (2006), astro-ph/0511573; D. Parkinson, P. Mukherjee, and A. R. Liddle, *Phys. Rev. D* **73**, 123523 (2006), astro-ph/0605003; M. Bridges, A. N. Lasenby, and M. P. Hobson, astro-ph/0607404; R. Trotta, astro-ph/0607496; R. Trotta, *Mon. Not. Roy. Astron. Soc.* **375**, L26 (2007), astro-ph/0608116; M. Kunz, B. A. Bassett, and R. Hlozek, astro-ph/0611004; P. Serra, A. F. Heavens, and A. Melchiorri, astro-ph/0701338; N. Bevis, M. Hindmarsh, M. Kunz, and J. Urrestilla, astro-ph/0702223; A. Niarchou and A. Jaffe, astro-ph/0702436; A. F. Heavens, T. D. Kitchling, and L. Verde, astro-ph/0703191
 - [7] B. A. Bassett, P. S. Corasaniti and M. Kunz, *Astrophys. J. Lett.* **617**, L1 (2004), astro-ph/0407364.
 - [8] P. Mukherjee, D. Parkinson, P. S. Corasaniti, A. R. Liddle, and M. Kunz, *Mon. Not. Roy. Astron. Soc.* **369**, 1725 (2006), astro-ph/0512484.
 - [9] M. Kunz, R. Trotta and D. Parkinson, *Phys. Rev. D* **74** 023503 (2006), astro-ph/0602378.
 - [10] C. Pahud, A. R. Liddle, P. Mukherjee, and D. Parkinson, *Phys. Rev. D* **73**, 083523 (2006), astro-ph/0605004.
 - [11] A. R. Liddle, P. Mukherjee, D. Parkinson, and Y. Wang, *Phys. Rev. D* **74**, 123506 (2006), astro-ph/0610126.
 - [12] R. Trotta, astro-ph/0703063.
 - [13] G. Schwarz, *Ann. Statist.* **5**, 461 (1978).
 - [14] A. R. Liddle, *Mon. Not. Roy. Astron. Soc.* **351**, L49, astro-ph/0401198; A. R. Liddle, astro-ph/0701113.
 - [15] D. Lindley, *Biometrika* **44**, 187 (1957).
 - [16] L. Lyons, ‘A particle physicist’s perspective on astro-statistics’, to appear in *Proceedings of the conference ‘Statistical Challenges in Modern Astronomy IV’*, State College, Pennsylvania, USA, June 2006.